

UiO : USIT

HPC - Best Practices

Best Practice Guide Modern Processors

Ole Widar Saastad, University of Oslo, Norway

Kristina Kapanova, NCSA, Bulgaria

Stoyan Markov, NCSA, Bulgaria

Cristian Morales, BSC, Spain

Anastasiia Shamakina, HLRS, Germany

Nick Johnson, EPCC, United Kingdom

Ezhilmathi Krishnasamy, University of Luxembourg, Luxembourg

Sebastien Varrette, University of Luxembourg, Luxembourg

Hayk Shoukourian (Editor), LRZ, Germany

Updated 5-5-2021



Ole W. Saastad, ITF/FI

15.03.2022

Best Practices HPC

First order approximation : «I've tested it and it works»

Higher orders : How to improve, efficiency, performance, experience, resource usage, good behavior, sharing experience, help onboarding¹

1:



Know your workload

Understanding what you run on the system is the key factor.

Remember you are not alone

HPC systems are shared by users who are all eager to get their work done!

Don't waste resources

**Is Parallel Programming Hard,
and, if so,
What Can You Do About It?**

Edited by Paul E. McKenney



What information do I need ?

How do I get this data ?

Gathering data about my work

Applications

- Single core / multicore ?
- Shared memory / Distributed memory ?

Gathering data about my work

The most important :

- Memory footprint
- Cores/processors - Scaling
- Input & output
- Run time

Applications

Single node application - single or multicore - Shared memory

Keywords : OpenMP, multicore, shared memory, threaded,
pthread, POSIX thread, NCPU

```
ldd ./pi.x | grep -i omp  
ldd ./pi.x | grep -i thread  
strings ./pi.x | grep OMP  
strings ./pi.x | grep thread
```

Applications

Multiple node application - single or multicore - Distributed memory

Keywords : MPI, distributed memory

```
ldd pi.x | grep mpi  
strings pi.x | grep mpi
```


Running applications

Memory, cores and run time

- Manual
- Other users
- Guess
- Trial and error

SLURM test jobs, exploring parameter space

Running applications

Interactively ?

Very short jobs (less than 5 min wall time) can be tested on front end

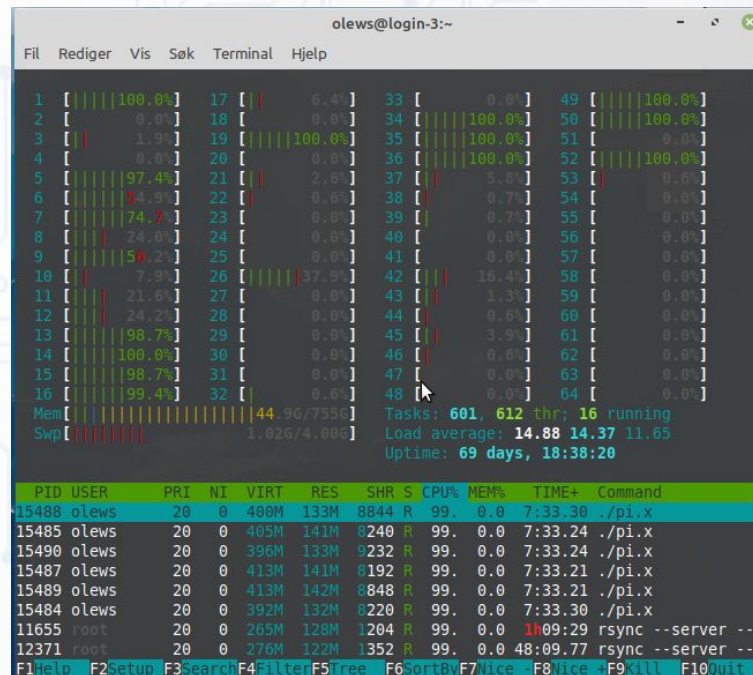
Reserve a node using SLURM, run interactively and explore.

Running applications

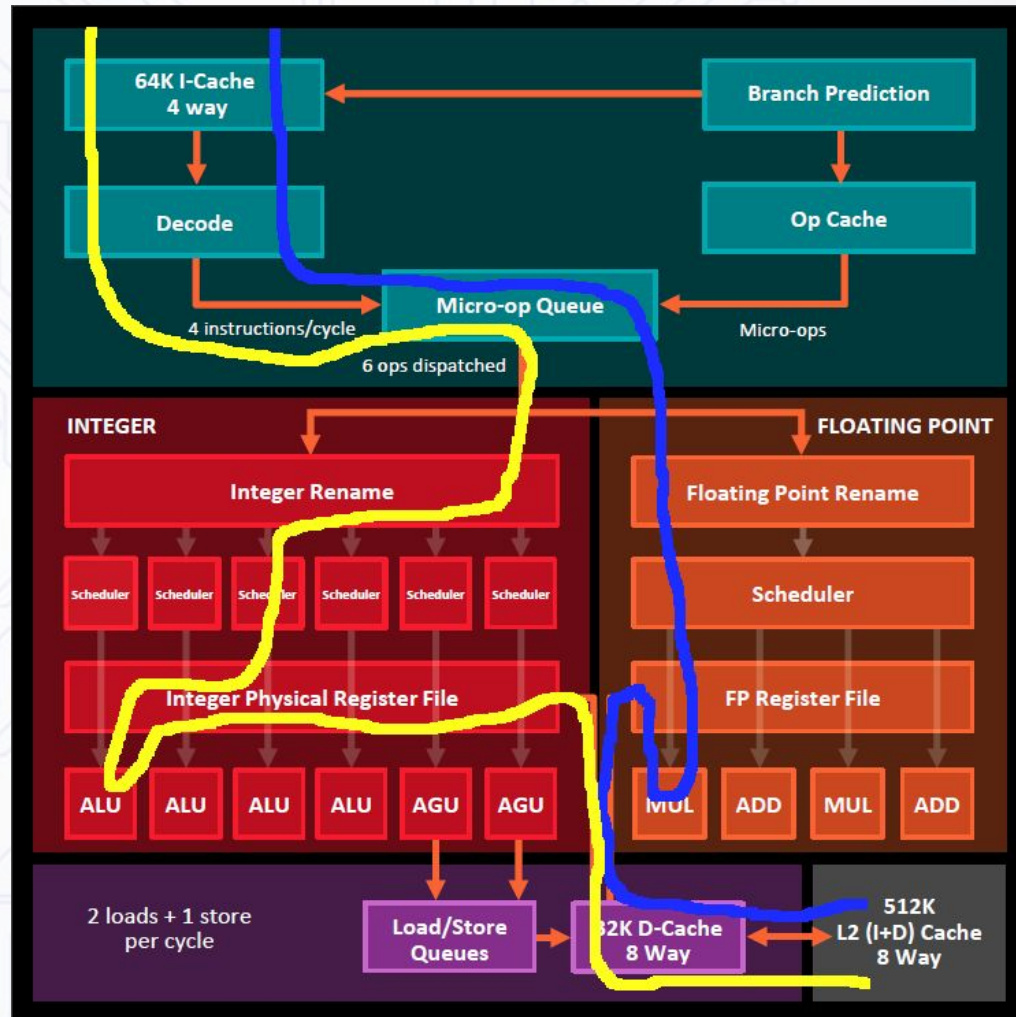
Interactive tools providing memory footprint and cores

- top
- htop

Review the man pages



A core running two instruction streams
(one integer + one floating point), two threads.



Running applications

SLURM test jobs, exploring parameter space

```
#SBATCH --nodes=2
```

```
#SBATCH --ntasks-per-node=4
```

```
#SBATCH --cpus-per-task=1
```

```
#SBATCH --time=00:10:00
```

```
#SBATCH --mem-per-cpu=2000M
```

This is a 5 dimensional space, only educated guesses are possible.

Running applications

SLURM test jobs, valuable feedback in logs:

Task and CPU usage stats:

JobID	JobName	AllocCPUS	NTasks	MinCPU	MinCPUTask	AveCPUElapsed	Exit Code
5312354	pi	8				00:01:02	0:0
5312354.bat+	batch	4	1	00:03:53	0	00:03:53	00:01:02 0:0
5312354.ext+	extern	8	2	00:00:00	0	00:00:00	00:01:02 0:0

Memory usage stats:

JobID	MaxRSS	MaxRSSTask	AveRSS	MaxPages	MaxPagesTask	AvePages
5312354						
5312354.bat+	1113772K	0	1113772K	0	0	0
5312354.ext+	0	0	0	0	0	0

Input and output - storage

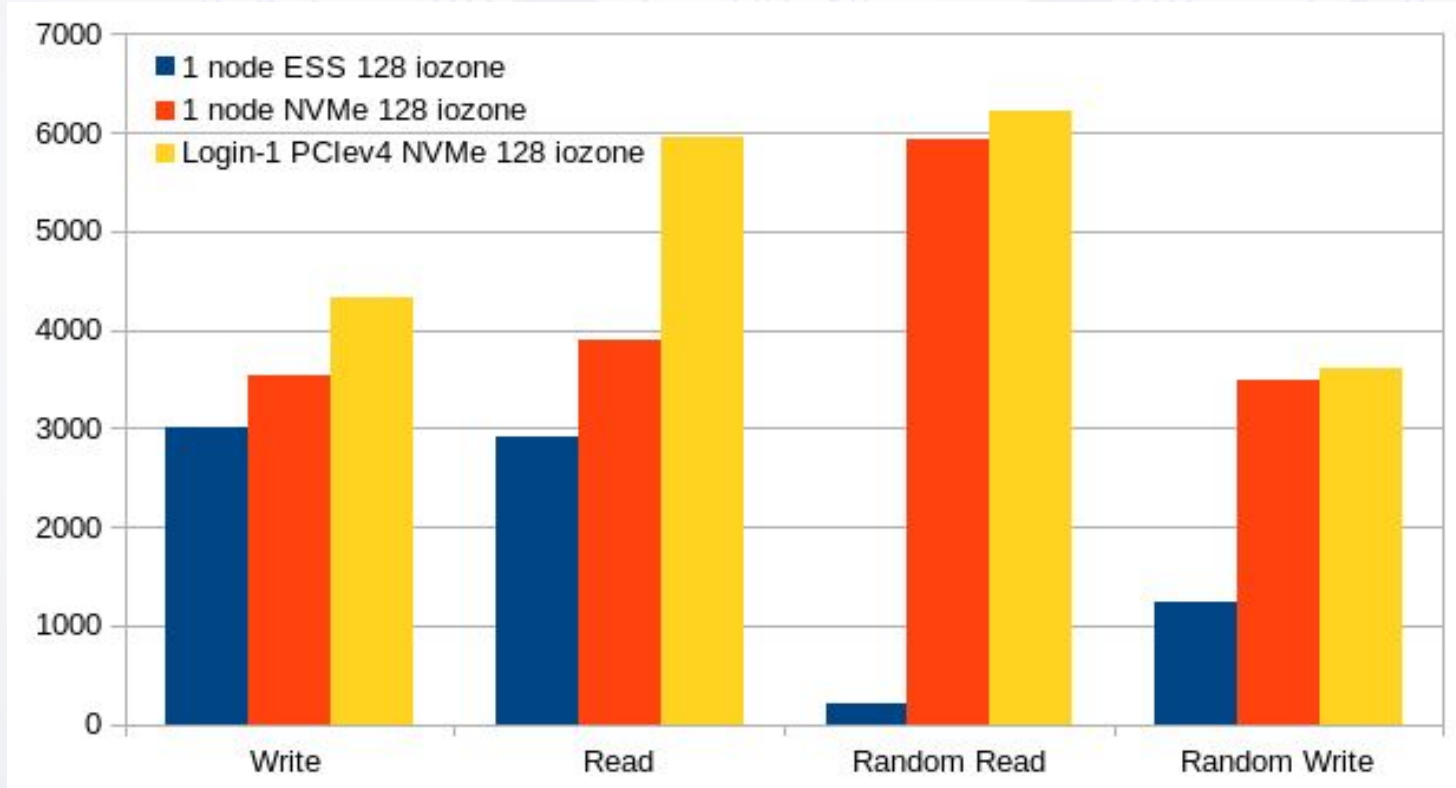
SLURM provide stats for IO

dd if=/dev/zero of=\$SCRATCH/dd.tmp bs=1024 count=10000000


Disk usage stats:

JobID	MaxDiskRead	MaxDiskReadTask	AveDiskRead	MaxDiskWrite	MaxDiskWriteTask	AveDiskWrite
5312653						
5312653.bat+	9765.78M	0	9765.78M	9765.63M	0	9765.63M
5312653.ext+	0.00M	0	0.00M	0	0	0

Input and output - storage



Scaling



Amdahl's law

$$Speedup = \frac{1}{(1 - p) + p/N}$$

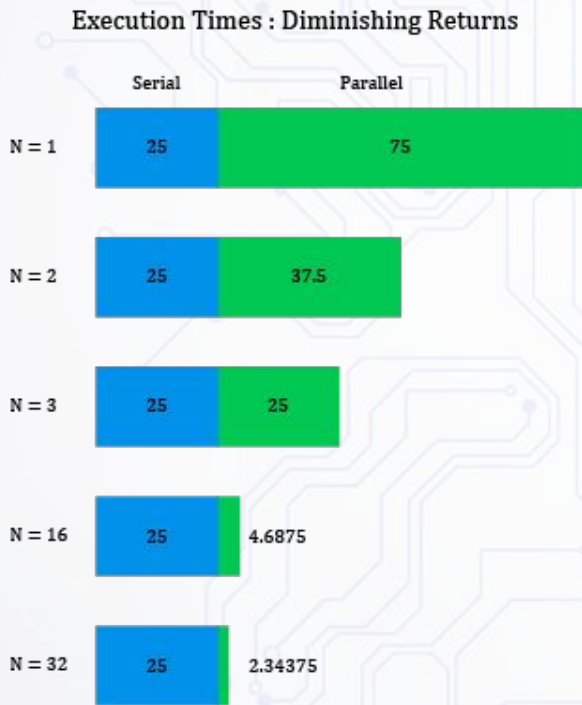
Suppose a car is traveling between two cities 60 miles apart, and has already spent one hour traveling half the distance at 30 mph. No matter how fast you drive the last half, it is impossible to achieve 90 mph average before reaching the second city. Since it has already taken you 1 hour and you only have a distance of 60 miles total; going infinitely fast you would only achieve 60 mph.

Gustavssons' law

$$\begin{aligned} S &= s + p \times N \\ &= s + (1 - s) \times N \\ &= N + (1 - N) \times s \end{aligned}$$

Suppose a car has already been traveling for some time at less than 90mph. Given enough time and distance to travel, the car's average speed can always eventually reach 90mph, no matter how long or how slowly it has already traveled. For example, if the car spent one hour at 30 mph, it could achieve this by driving at 120 mph for two additional hours, or at 150 mph for an hour, and so on.

Amdahl's law & Gustavssons' law



</projects/ec34/inf9380/Exercises-23.Mar/>